

FURONG JIA

jiafurong2002@gmail.com | flora-jia-jfr.github.io/furongjia.github.io/

Durham, North Carolina, United States

RESEARCH STATEMENT

I am broadly interested in making AI systems, especially large language models, **reliable, interpretable, and robust** by understanding *how data shapes model behavior at training and inference time*. My interests include methods that leverage properties of diverse data (text, time-series, multimodal) to make models adaptable to changes that arise during both training and deployment. A parallel interest is applying these approaches in clinical contexts, where robustness, generalizability, and responsible deployment are critical.

EDUCATION

• Duke University

Doctor of Philosophy in Computer Science

◦ Advisor: Monica Agrawal

Aug 2024 - Present

Durham, NC

GPA: 4.0/4.0

• University of Southern California

Bachelor of Science in Computer Science & Bachelor of Science in Applied and Computational Mathematics

◦ W.V.T. Rusch Engineering Honors Program, Viterbi Grand Challenges Scholar

Aug 2020 - May 2024

Los Angeles, CA

GPA: 3.96/4.0

PUBLICATIONS

Conference:

[1] Interpreting Dataset Shift in Clinical Notes

ML4H 2025 (Machine Learning for Health Symposium)

Shariar Vaez-Ghaemi*, Furong Jia*, Monica Agrawal.

[2] Diagnosing our datasets: How does my language model understand clinical text?

CHIL 2025 (7th Annual Conference on Health, Inference, and Learning)

Furong Jia, David Sontag, Monica Agrawal.

[3] GPT4MTS: Prompt-based Large Language Model for Multimodal Time-series Forecasting.

EAAI-24 (The 14th Symposium on Educational Advances in Artificial Intelligence)

Furong Jia, Kevin Wang, Yixiang Zheng, Defu Cao, Yan Liu. (2024).

[4] TEMPO: Prompt-based Generative Pre-trained Transformer for Time Series Forecasting.

ICLR 2024 (The Twelfth International Conference on Learning Representations)

Defu Cao, Furong Jia, Sercan O Arik, Tomas Pfister, Yixiang Zheng, Wen Ye, Yan Liu.

[5] I2I: Initializing Adapters with Improvised Knowledge.

CoLLAs 2023 (Second Conference on Lifelong Learning Agents)

Tejas Srinivasan, Furong Jia, Mohammad Rostami, Jesse Thomason.

Workshop/Findings:

[1] Counting Clues: A Lightweight Probabilistic Baseline Can Match an LLM

ML4H 2025 Findings (Machine Learning for Health Symposium)

Furong Jia*, Yuan Pu*, Finn Guo, Monica Agrawal.

[2] What Patients Really Ask: Exploring the Effect of False Assumptions in Patient Information Seeking

ML4H 2025 Findings (Machine Learning for Health Symposium)

Raymond M Xiong, Furong Jia, Lionel Wong, Monica Agrawal.

Under Review:

[1] Batch-of-Thought: Cross-Instance Learning for Enhanced LLM Reasoning

Xuan Yang, Furong Jia, Roy Xie, Xi Xiong, Jian Li, Monica Agrawal

RESEARCH INTEREST

Natural Language Processing, Data-centric Machine Learning, Healthcare, Time Series, Multimodal Machine Learning

RESEARCH EXPERIENCE

• DukeNLP, Duke University

Advisor: Monica Agrawal

Aug 2024 - Present
Durham, NC

◦ Understanding LLM dynamics as a function of pretraining data

- * We examined the relationship of clinical information (clinical jargon understanding and unsupported medical claims) between the occurrence of keyword pairs in pretraining corpora and model performance. For clinical jargon understanding, we also curated a benchmark *Medlingo* using clinical abbreviations from real-world clinical notes. We further analyzed the sources of clinical information in pretraining datasets to guide future corpus design. [4].
- * We tested whether success on clinical multiple-choice diagnostic questions reflects probabilistic inference by introducing the *Frequency-Based Diagnostic Ranker (FBDR)*, a simple Naive Bayes method that extracts clinical concepts and scores diagnoses from co-occurrence statistics in the same pretraining corpora as the evaluated models. FBDR achieved accuracy comparable to a 7B LLM using the same pretraining data. The two systems made different correct predictions, which indicates complementary strengths. These results suggest that model performance involves mechanisms beyond frequency aggregation, while a transparent, low complexity expert system style baseline still captures a meaningful share of the benchmark signal and is a useful target for hybrid approaches [2].
- * Real-world decisions often rely on calibrated likelihoods through probabilistic reasoning rather than certainties (i.e. clinical diagnostics), we are motivated to understand where LLMs acquire such probabilistic knowledge during pretraining. Our ongoing work investigates how pretraining data equips LLMs with probabilistic knowledge through correlated entities and to what extent they can internalize abstract dependency structures through continual pretraining on synthetic data.

◦ Interpreting Dataset Shift in Clinical Notes

- * Distribution shift degrades ML performance, especially in clinical text. Therefore, actionability requires not just detection but *explanation*. We establish an extensible benchmark suite that induces synthetic distribution shifts using real clinical notes and develop two methods for assessing generated shift explanations. We further introduce *SIRENs*, a general-domain end-to-end approach that explains distributional differences by selecting representative notes from each. SIRENs reliably recover salient binary shifts with comparatively lower performance on subtle continuous changes, showing a gap to a ground-truth oracle and suggesting room for improvement in future methods. [1]

• Melady Lab, University of Southern California

May 2023 - June 2024
Los Angeles, CA

Advisor: Yan Liu

◦ Large Language Models for Time Series Forecasting

- * We introduced *GPT4MTS*, a prompt-based framework that fuses numeric time series with aligned textual context. We also built a GDELT-derived multimodal news impact dataset and showed consistent forecasting gains over strong unimodal baselines, demonstrating effective multimodal fusion and the value of extra-textual information. [5]
- * We developed *TEMPO*, a GPT architecture-based time series framework that combines seasonal trend decomposition with adaptive, selection-based prompting for non-stationary distribution shift. It achieves zero-shot state-of-the-art across diverse and unseen (including multi-modal) benchmarks, indicating strong potential as a foundation model for time series forecasting. [6]

• GLAMOR Lab, University of Southern California

Aug 2022 - Present
Los Angeles, CA

Advisor: Jesse Thomason

◦ Initializing Adapters with Improvised Knowledge for Multimodal Continual Learning

- * We applied parameter-efficient knowledge transfer using adapters in multimodal models for visual question-answering tasks. We implemented and trained model architectures on multiple vision-language tasks, including GQA, COCO-QA, and AQUA, to enhance forward transfer and reduce knowledge forgetting. [7]

WORK EXPERIENCE

• Galileo Financial Technologies, SoFi

May 2022 - Aug. 2022

US Remote

Software Engineering Intern

- Independently designed, developed, and tested micro-services in Flask to support authentication and retrieve account information from the database, and generated scripts for Continuous Integration through Gitlab.
- Independently designed, developed, and tested an account-information-component that creates web components to display an account information card using the lit-element framework in HTML/CSS and Typescript.

HONORS AND AWARDS

- **Albert Dorman Future Leader Award**, University of Southern California 2024
- **Student Recognition Awards**, University of Southern California 2024
- **Provost's Research Fellowship**, University of Southern California Fall 2023
- **CURVE Research Fellowship**, University of Southern California Fall 2022, Spring 2023
- **Academic Achievement Award Scholarship**, University of Southern California 2021
- **ABC Innovation Prize**, University of Southern California 2021

SERVICE

- Teaching Assistant:
 - * Applied Machine Learning, *Duke University* Spring 2025
 - * Algorithms and Theory of Computing, *University of Southern California* Spring 2022
 - * Web Publishing and Front-end Development, *University of Southern California* Fall 2021 - Spring 2024
- Reviewer:
 - * ACL ARR 2024 (June, October), ACL ARR 2025 (February, May)
 - * NeurIPS'25
 - * ML4H'25